

R Workshop: ANOVA, ANCOVA & Factorial ANOVA

Brier Gallihugh, M.S.

2023-06-18

Table of contents

Introduction to ANOVA	2
Running ANOVA	2
Assumptions of ANOVA	3
Model Normality	3
Homogeneity of Variance	6
Introduction to ANCOVA	7
Running ANCOVA	7
Assumptions of ANCOVA	10
Predictor x Covariate Independence	10
Homogeneity of Regression Slopes	10
Residual Normality (DV ~ IV)	11
Homogeneity of Variance	12
Linearity of CV & DV	13
Introduction to Factorial ANOVA	14
Running Factorial ANOVA	14
Assumptions of Factorial ANOVA	16
Linearity of IV to DV (Regression)	16
Homogeneity of Variance (Regression)	17
Normality of Residuals	20

Introduction to ANOVA

Running ANOVA

```
library(car)
library(psych)
library(multcomp)
library(effects)
library(tidyverse)
library(sjstats)

data <- dplyr::starwars

data <- data %>%
  dplyr::select(height,mass,hair_color,species,sex) %>%
  na.omit() %>%
  filter(species == "Gungan" | species == "Human" | species == "Wookiee")①

height_species_aov <- aov(height ~ species, data = data) ②
summary(height_species_aov) ③

# Post Hoc Test (Tukey)
TukeyHSD(height_species_aov) ④
# Bonferroni Adjustment vs Tukey
pairwise.t.test(data$height, data$species, p.adjust.method = "bonferroni")⑤
```

- ① Here we are using the `filter()` function to filter for “Gungan”, “Human” and “Wookiee”
- ② This is the `aov()` function. Here we are specifying the ANOVA formula.
- ③ The `summary()` function will give you the output of the ANOVA
- ④ Because the `species` variable has more than 2 conditions, we might want to figure out which comparisons are statistically different from each other. The `TukeyHSD()` function allows us to perform a post hoc Tukey test
- ⑤ Maybe Tukey isn’t your favorite correction. Another route might be to use a pairwise t test that corrects for multiple comparisons using a Bonferroni correction. You can do that with the `pairwise.t.test()` function.

```
      Df Sum Sq Mean Sq F value    Pr(>F)
species  2   6118    3059    22.5 3.86e-06 ***
Residuals 23   3127     136
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
 Tukey multiple comparisons of means
 95% family-wise confidence level

Fit: aov(formula = height ~ species, data = data)

```
$species
```

	diff	lwr	upr	p adj
Human-Gungan	-30.45455	-52.022229	-8.886862	0.0048319
Wookiee-Gungan	21.00000	-8.202782	50.202782	0.1915885
Wookiee-Human	51.45455	29.886862	73.022229	0.0000125

Pairwise comparisons using t tests with pooled SD

data: data\$height and data\$species

	Gungan	Human
Human	0.0053	-
Wookiee	0.2545	1.3e-05

P value adjustment method: bonferroni

Assumptions of ANOVA

As ANOVA is just a special case of linear regression, the same assumptions exist for ANOVA that exist for linear regression. We will test each assumption below. The code will look awfully similar to the code used for linear regression previously

Model Normality

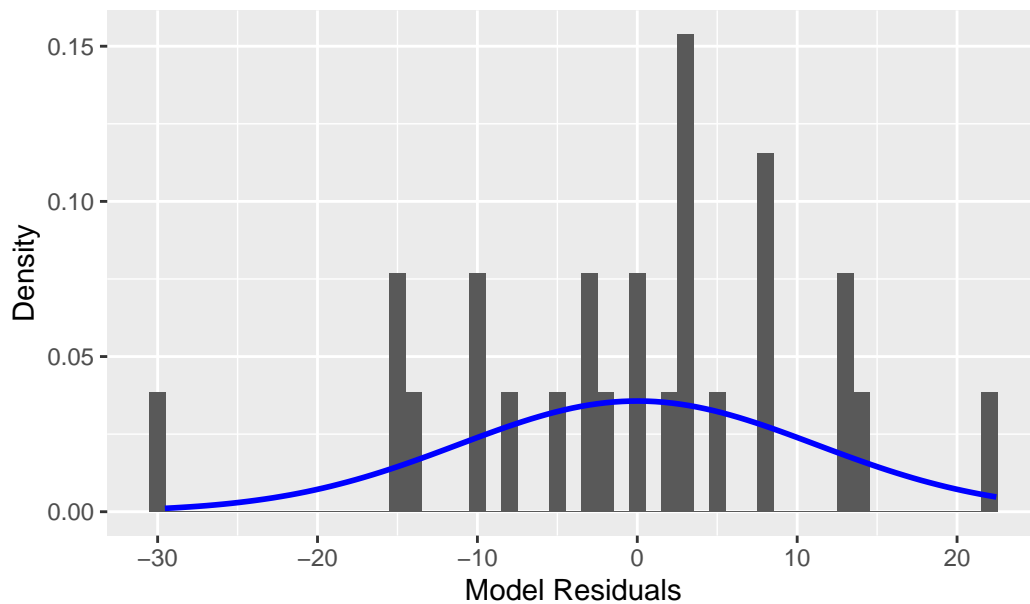
Graphical

```
density_plot <- data %>% ggplot(aes(x = height_species_aov$residuals)) +
  geom_histogram(aes(y= after_stat(density)),binwidth = 1) +
  stat_function(fun = dnorm,
               args = list(mean = mean(height_species_aov$residuals),
                           sd = sd(height_species_aov$residuals)),
               col = "blue",
               linewidth = 1) +
  labs(title = "Figure 1. Histogram of Model Residual Scores",
```

```
x = "Model Residuals",  
y = "Density")
```

```
density_plot
```

Figure 1. Histogram of Model Residual Scores




```
# OR
```

```
qqnorm(height_species_aov$residuals)  
qqline(height_species_aov$residuals, col = "blue")
```

①

Normal Q-Q Plot



1. When running the `qqnorm()` and `qqline()` functions at the same time, you will get a qq-plot which will graphically assess the normality of the argument given (in this case the model residuals). The `col` argument simply tells R which color to make the reference line.

Statistical

```
psych::describe(height_species_aov$residuals)
shapiro.test(height_species_aov$residuals)
```

①

① Another way to assess normality is to look at the skew and kurtosis of the data. Typically skew and kurtosis between -1 and 1 are acceptable. We can see this using the `describe()` function within the `psych` package.

```
vars  n mean    sd median trimmed  mad   min   max range  skew kurtosis
X1    1  26    0 11.18   1.45   0.35 10.38 -29.55 22.45   52 -0.44    0.16
se
X1 2.19
```

Shapiro-Wilk normality test

```
data: height_species_aov$residuals
W = 0.97347, p-value = 0.7143
```

Homogeneity of Variance


As with regression, homogeneity of variance is also important for ANOVA. We will graphically and statistically assess this assumption below.

Graphical

```
boxplot(data$height ~ data$species)
```

①

- ① The `boxplot()` function will give us a visual representation of the DV at each level of the IV. We are looking for boxplots that are roughly the same height across each group.



Statistical

Statistically, we can assess homogeneity of variance using Levene's Test. Here we want the results to not be statistically significant.

```
leveneTest(data$height ~ data$species)
```

①

- ① The `leveneTest()` function from the `car` package takes a DV and IV as arguments.

Levene's Test for Homogeneity of Variance (center = median)

	Df	F value	Pr(>F)
group	2	1.168	0.3288
	23		

Introduction to ANCOVA

Running ANCOVA

ANCOVA refers to any ANOVA model with 2 or more parameters. These models require similar assumptions to ANOVA with a couple of additional ones. Below we will see how to run an ANCOVA and then test its assumptions.

```
data <- data %>% mutate(species = as.factor(species))
# Type I SS
ancova <- aov(height ~ species + mass, data = data)
summary(ancova) ①
# Reversed Order
ancova2 <- aov(height ~ mass + species, data = data)
summary(ancova2) ②

car::Anova(ancova2, type = "III") ③

# Std Means
summary_data <- data %>%
  dplyr::group_by(species) %>%
  dplyr::summarise(n = n(),
                   mean = mean(height))

# Adjust means for covariate effect
adjustedMeans<- effect("species", ancova, se=TRUE) ④
summary(adjustedMeans)
adjustedMeans$se

# Post Hoc Tests

posthoc <- multcomp::glht(ancova, linfct = multcomp::mcp(species = "Tukey"))⑤
summary(posthoc)
confint(posthoc)

# Effect size
sjstats::anova_stats(car::Anova(ancova,type = "III")) ⑥
```

- ① This displays our ANCOVA output with species and then mass.
- ② This displays our ANCOVA output with mass and then species. We can see that 1 and 2 are different. This is because R by default does a sequential ANCOVA so order matters.

- ③ To get around this, we can use the `Anova()` function in the `car` package to specify that we want Type III sums of squares. This will then generate an ANCOVA where the order of predictors doesn't matter
- ④ When having covariates, one might wish to report adjusted means. We can do that using the `effect()` function from the `effects` package.
- ⑤ To run a post hoc analysis, we need to use the `ghlt()` function from the `multcomp` package. There are additional corrections you can do outside of Tukey. To investigate, please see the documentation.
- ⑥ To get basic ANOVA effect size statistics, we can use the `anova_stats()` function from the `sjstats` package.

```

      Df Sum Sq Mean Sq F value    Pr(>F)
species     2   6118   3059.0    37.27 8.61e-08 ***
mass        1   1322   1321.8    16.10 0.000584 ***
Residuals   22   1806     82.1
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      Df Sum Sq Mean Sq F value    Pr(>F)
mass        1   3737   3737    45.53 8.81e-07 ***
species     2   3702   1851    22.55 4.70e-06 ***
Residuals   22   1806     82
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Anova Table (Type III tests)

Response: height
      Sum Sq Df F value    Pr(>F)
(Intercept) 28097.6  1 342.337 6.734e-15 ***
mass        1321.8  1  16.104 0.0005843 ***
species     3702.4  2  22.555 4.699e-06 ***
Residuals   1805.7 22
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

species effect
species
  Gungan   Human  Wookiee
214.4986 180.5408 215.5526

Lower 95 Percent Confidence Limits
species
  Gungan   Human  Wookiee

```


201.0113 176.5022 200.0533

Upper 95 Percent Confidence Limits

species

Gungan Human Wookiee
227.9859 184.5794 231.0520
[1] 6.503427 1.947367 7.473624

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts

Fit: aov(formula = height ~ species + mass, data = data)

Linear Hypotheses:

	Estimate	Std. Error	t value	Pr(> t)
Human - Gungan == 0	-33.958	6.748	-5.033	<0.001 ***
Wookiee - Gungan == 0	1.054	10.333	0.102	0.994
Wookiee - Human == 0	35.012	7.846	4.462	<0.001 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
(Adjusted p values reported -- single-step method)

Simultaneous Confidence Intervals

Multiple Comparisons of Means: Tukey Contrasts

Fit: aov(formula = height ~ species + mass, data = data)

Quantile = 2.4824

95% family-wise confidence level

Linear Hypotheses:

	Estimate	lwr	upr
Human - Gungan == 0	-33.9578	-50.7079	-17.2076
Wookiee - Gungan == 0	1.0541	-24.5973	26.7054
Wookiee - Human == 0	35.0118	15.5356	54.4880

term	sumsq	meansq	df	statistic	p.value	etasq	partial.etasq	omegasq
------	-------	--------	----	-----------	---------	-------	---------------	---------

species		3702.445		1851.223		2		22.555		< .001		0.542		0.672		0.512
mass		1321.785		1321.785		1		16.104		0.001		0.194		0.423		0.179
Residuals		1805.669		82.076		22										

Assumptions of ANCOVA

ANCOVA requires the same assumptions as ANOVA with two additional assumptions: 1) That the predictor and covariate are independent 2) The regression slopes are homogeneous

We will test each of the traditional parametric tests as well as the two additional assumptions below.

Predictor x Covariate Independence

Statistical

```
predictor_assumption <- aov(mass ~ species, data = data) ①
summary(predictor_assumption)
```

- ① To assess this statistically, one needs to run an ANOVA looking at each predictor with each additional parameter. We do not want these to be statistically significant.

```

              Df Sum Sq Mean Sq F value Pr(>F)
species        2   3390   1695.1    4.694 0.0195 *
Residuals     23   8306    361.1
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Tip

Here the statistical test for this assumption is statistically significant (which is bad). It actually means we can't do this (unless we have random assignment)

Homogeneity of Regression Slopes

Statistical

```
regression_slope_assumption <- aov(height ~ species*mass, data = data)
car::Anova(regression_slope_assumption, type = "III") ①
```

- ① To test this assumption, one just needs to test the interaction effects within the model. The code for this is shown here.

Anova Table (Type III tests)

Response: height

	Sum Sq	Df	F value	Pr(>F)
(Intercept)	149.72	1	1.9153	0.18163
species	155.01	2	0.9914	0.38856
mass	392.00	1	5.0144	0.03666 *
species:mass	242.18	2	1.5490	0.23689
Residuals	1563.48	20		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Tip

For the above example, the interaction is NOT significant so assumption is met

Residual Normality (DV ~ IV)

Residual normality is an assumption shared with standard regression. We can assess it the same way we did previously when looking at simple ANOVA. The same is true for the homogeneity of variance assumption.

Statistical

```
height_mass_resid <- scale(residuals(aov(height ~ mass, data = data))) ①
psych::describe(height_mass_resid) ②
shapiro.test(height_mass_resid) ③
```

- ① The `scale()` function will scale the residuals of the ANOVA. The `residuals()` function pulls out the residuals within the ANOVA
- ② Statistical summary of the residuals using the `describe()` function in the `psych` package
- ③ Statistical test (Shapiro Wilks) of the residuals using the `shapiro.test()` function

```
vars  n mean sd median trimmed  mad   min max range skew kurtosis  se
X1    1 26   0  1      0   -0.06 0.82 -1.85 2.7  4.55 0.71    0.44 0.2
```

Shapiro-Wilk normality test


```
data: height_mass_resid
W = 0.95143, p-value = 0.2505
```

Graphical

```
hist(height_mass_resid, col = 'beige',
      main="", xlab = "ANCOVA Residuals (z-score)",
      probability = TRUE)
curve(dnorm(x, mean = mean(height_mass_resid),
            sd = sd(height_mass_resid)),
      add = TRUE, lwd = 2, col = 'blue')
```

①

① A graphical histogram (with normal distribution overlay) of the model residuals



Homogeneity of Variance

Statistical

```
# Levene's Test to Assess Equal Variance for Species
car::leveneTest(data$height ~ data$species)
```

①

- ① A statistical test (Levene's test) of the homogeneity assumption using the `leveneTest()` function in the `car` package.


Levene's Test for Homogeneity of Variance (center = median)

```
Df F value Pr(>F)
group 2    1.168 0.3288
      23
```

Graphical

```
# Boxplot Height by Species
boxplot(height ~ species, data=data, main="Height Variance by Species ", ①
        xlab="Species", ylab="Height")
```

- ① A visual representation of the homogeneity of variance assumption using a boxplot



Linearity of CV & DV


The last assumption is that the CV and DV are linearly related. We can test this using the code below graphically.

Graphical

```
plot(lm(data$height ~ data$mass, data = data),  
     pch = 16, bty = 'l', 2)
```

①

① Visual representation of the CV to DV linearity assumption



Introduction to Factorial ANOVA

Running Factorial ANOVA

```
library(tidyverse)
factorial_data <- starwars %>%
  select(mass, homeworld, species) %>%
  mutate(homeworld = as.factor(homeworld),
         species = as.factor(species)) %>%
  filter(homeworld == "Tatooine" | homeworld == "Naboo") %>%
  filter(species == "Human" | species == "Droid") %>% na.omit()
# Modeling As Factorial ANOVA
aov_factorial <- aov(mass ~ species*homeworld, data = factorial_data)
Anova(aov_factorial, type = "III")

# Modeling As A Regression (For SS Analyses)
reg_fanova <- lm(mass ~ homeworld*species, data = factorial_data)
summary(reg_fanova) # <10>
```

①

②

③

④

⑤

⑥

⑦

⑧

⑨

- ① Start with the `starwars` data set
- ② Use the `select()` function to pull out mass, homeworld and species variables
- ③ Call the `mutate()` function to format the homeworld variable using the `as.factor()` function.
- ④ Format the species variables as a factor using the `as.factor()` function.
- ⑤ Use the `filter()` function to filter for observations with homeworld = "Tatooine" OR "Naboo"
- ⑥ Use the `filter()` function to filter for observations with species = "Human" OR "Droid"
- ⑦ Create an ANOVA object using the `aov()` function
- ⑧ Use the `Anova()` function from the `car()` package to specify Type III Sums of Squares
- ⑨ Create an Regression object using the `lm()` function
- ⑩ Show output of the regression object using the `summary()` function.

Anova Table (Type III tests)

Response: mass

	Sum Sq	Df	F value	Pr(>F)
(Intercept)	1024.0	1	1.6199	0.2388
species	990.1	1	1.5662	0.2461
homeworld	308.2	1	0.4875	0.5048
species:homeworld	19.0	1	0.0301	0.8666
Residuals	5057.2	8		

Call:

```
lm(formula = mass ~ homeworld * species, data = factorial_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-23.33	-19.50	-6.00	17.88	40.00

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	32.000	25.143	1.273	0.239
homeworldTatooine	21.500	30.793	0.698	0.505
speciesHuman	36.333	29.032	1.251	0.246
homeworldTatooine:speciesHuman	6.167	35.557	0.173	0.867

Residual standard error: 25.14 on 8 degrees of freedom

Multiple R-squared: 0.5219, Adjusted R-squared: 0.3426

F-statistic: 2.91 on 3 and 8 DF, p-value: 0.1009

Assumptions of Factorial ANOVA


Linearity of IV to DV (Regression)

Graphical

```
plot(lm(mass ~ species, data = factorial_data), 2)
```


①

② Plot the residuals of the model using the `plot()` and `lm()` functions for mass on homeworld



```
plot(lm(mass ~ homeworld, data = factorial_data), 2)
```

②



Statistical

```
psych::describe(resid(aov_factorial))
shapiro.test(resid(aov_factorial))
```

①
②

- ① Statistical summary of the model residuals using the `describe()` function in the `psych` package
- ② Statistical test (Shapiro Wilk) of the residuals using the `resid()` and `shapiro.test()` functions

```
vars  n mean    sd median trimmed  mad   min max range skew kurtosis   se
X1    1 12   0 21.44    -6  -1.67 22.61 -23.33  40 63.33 0.43    -1.4 6.19
```

Shapiro-Wilk normality test

```
data:  resid(aov_factorial)
W = 0.90166, p-value = 0.1666
```


Homogeneity of Variance (Regression)

Graphical

```
boxplot(factorial_data$mass ~ factorial_data$species)
```

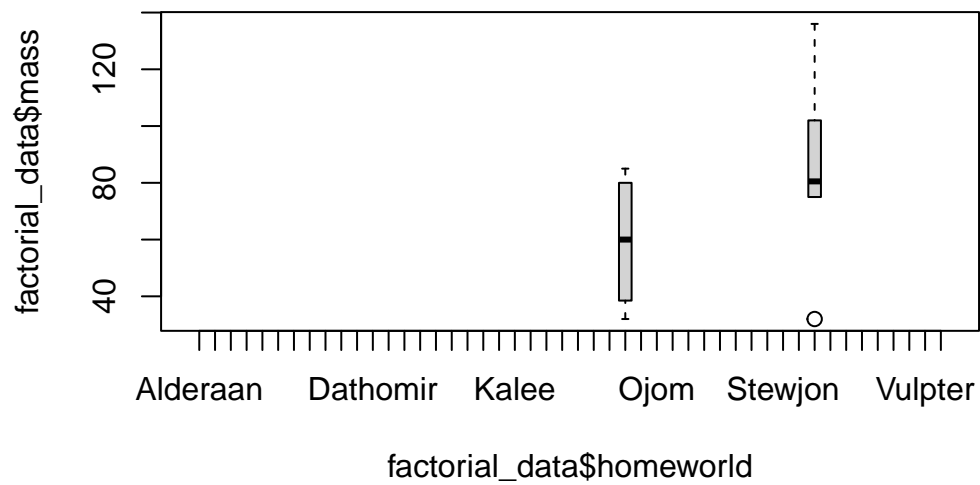
①

- ③ Fitted object of the ANCOVA
- ④ Plot of the residual x fitted data using the `plot()` function



```
boxplot(factorial_data$mass ~ factorial_data$homeworld)
```

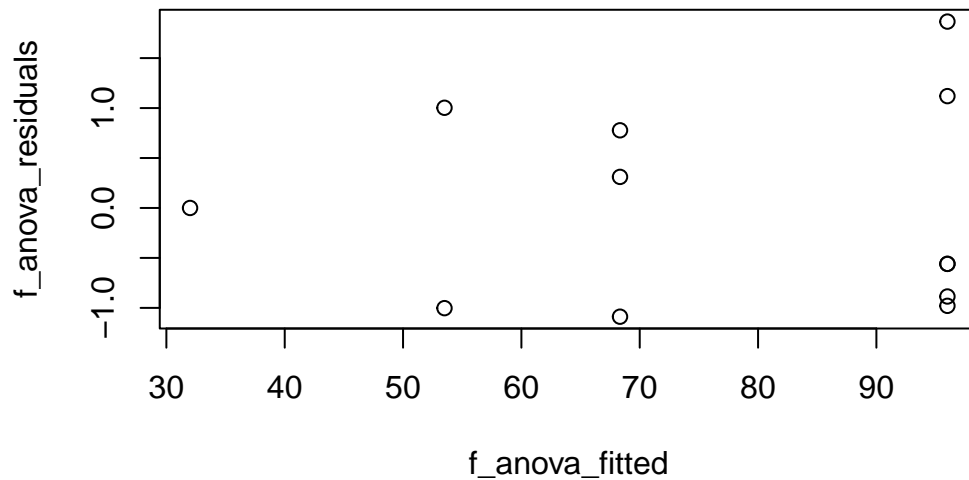
①



```
# Residuals
f_anova_residuals <- scale(residuals(lm(mass ~ species*homeworld, data = factorial_data)))
f_anova_fitted <- fitted(lm(mass ~ species*homeworld, data = factorial_data))

plot(f_anova_residuals ~ f_anova_fitted)
```

④



Statistical

```
leveneTest(factorial_data$mass ~ factorial_data$species) ①
```

```
leveneTest(factorial_data$mass ~ factorial_data$homeworld) ②
```

```
leveneTest(factorial_data$mass ~ interaction(factorial_data$homeworld, factorial_data$species)) ③
```

- ① Levene test of mass x species using `leveneTest()` function
- ② Levene test of mass x homeworld using `leveneTest()` function
- ③ Levene test of mass x species/homeworld interaction using `leveneTest()` function

Levene's Test for Homogeneity of Variance (center = median)

	Df	F value	Pr(>F)
group 1	1	0.0367	0.852
	10		

Levene's Test for Homogeneity of Variance (center = median)

	Df	F value	Pr(>F)
group 1	1	1e-04	0.9919
	10		

Levene's Test for Homogeneity of Variance (center = median)

	Df	F value	Pr(>F)
group 3	3	0.3306	0.8037
	8		


Normality of Residuals

Graphical

```
plot(aov_factorial, 2)
```

①

- ① Plot ANOVA residuals using `plot()` function



Statistical

```
fanova_residuals <- aov_factorial$residuals
```

①

```
shapiro.test(fanova_residuals)
```

②

```
psych::describe(fanova_residuals)
```

③

- ① Create residual data set
- ② Statistical test of residual normality using `shapiro.test()` function
- ③ Descriptive statistics of residuals using `describe()` function in the `psych` package

Shapiro-Wilk normality test

```
data: fanova_residuals  
W = 0.90166, p-value = 0.1666
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X1	1	12	0	21.44	-6	-1.67	22.61	-23.33	40	63.33	0.43	-1.4	6.19